

Using the PhysX engine for physics-based virtual surgery with force feedback

Anderson Maciel^{1*}

Tansel Halic²

Zhonghua Lu²

Luciana P. Nedel¹

Suvranu De²

¹*Instituto de Informática,
Universidade Federal do Rio Grande
do Sul, Porto Alegre, Brazil*

²*Department of Mechanical,
Aerospace and Nuclear Engineering,
Rensselaer Polytechnic Institute, Troy,
USA*

*Correspondence to:

Anderson Maciel, Instituto de
Informática, UFRGS, Campus do
Vale, Bloco IV, Av. Bento Gonçalves
9500, Caixa Postal 15.064, CEP
91501-970 Porto Alegre, RS, Brazil.
E-mail: amaciel@inf.ufrgs.br

Abstract

Background The development of modern surgical simulators is highly challenging, as they must support complex simulation environments. The demand for higher realism in such simulators has driven researchers to adopt physics-based models, which are computationally very demanding. This poses a major problem, since real-time interactions must permit graphical updates of 30 Hz and a much higher rate of 1 kHz for force feedback (haptics). Recently several physics engines have been developed which offer multi-physics simulation capabilities, including rigid and deformable bodies, cloth and fluids. While such physics engines provide unique opportunities for the development of surgical simulators, their higher latencies, compared to what is necessary for real-time graphics and haptics, offer significant barriers to their use in interactive simulation environments.

Methods In this work, we propose solutions to this problem and demonstrate how a multimodal surgical simulation environment may be developed based on NVIDIA's PhysX physics library. Hence, models that are undergoing relatively low-frequency updates in PhysX can exist in an environment that demands much higher frequency updates for haptics. We use a collision handling layer to interface between the physical response provided by PhysX and the haptic rendering device to provide both real-time tissue response and force feedback.

Results Our simulator integrates a bimanual haptic interface for force feedback and per-pixel shaders for graphics realism in real time. To demonstrate the effectiveness of our approach, we present the simulation of the laparoscopic adjustable gastric banding (LAGB) procedure as a case study.

Conclusions To develop complex and realistic surgical trainers with realistic organ geometries and tissue properties demands stable physics-based deformation methods, which are not always compatible with the interaction level required for such trainers. We have shown that combining different modelling strategies for behaviour, collision and graphics is possible and desirable. Such multimodal environments enable suitable rates to simulate the major steps of the LAGB procedure. Copyright © 2009 John Wiley & Sons, Ltd.

Keywords computer graphics; interaction techniques; physics-based simulation; minimally invasive surgery; user interfaces–haptic I/O

Introduction

Physics-based interactions are necessary to improve the realism of modern surgical simulators and allow the response to surgical interventions to be

Accepted: 6 April 2009

computed in a physiologically relevant manner. For real-time graphical rendering, an update rate of at least 30 Hz must be maintained, whereas for stable force feedback (haptics) a much higher update rate of the order of 1 kHz is necessary. The availability of new powerful graphics hardware is insufficient to meet the demands of such physics-based surgical simulators, which must now include interactions of surgical tools with multiple deformable objects, bleeding and smoke generation due to cautery procedures, as well as tissue approximation procedures such as suturing and stapling. Recently several physics engines have been developed that offer multi-physics simulation capabilities, including rigid and deformable bodies, cloth and fluids. While such physics engines provide unique opportunities for the development of surgical simulators, their higher latencies, compared to what is necessary for real-time graphics and haptics, offer significant barriers to their use in interactive simulation environments.

In this study, we propose solutions to this problem and demonstrate how a multimodal surgical simulation environment may be developed based on NVIDIA's PhysX physics library. PhysX, which can be accelerated using the physics processing unit (PPU) or CUDA-enabled GeForce graphics processing unit (GPU), provides an optimized set of methods for physics-based simulation. Our simulator integrates a bimanual haptic interface for force feedback and shaders for graphic realism in real time. A model-view-controller (MVC) architecture, taking advantage of the parallel computation of modern GPU and multi-core CPU, has been developed to integrate PhysX to exploit the hardware to the maximum. In this integration, heterogeneous models – triangle mesh, tetrahedron mesh, rigid, deformable and articulated – cohabit in the environment and allow equally heterogeneous three-dimensional (3D) interaction modes: line- and point-based collision detection, with or without force feedback, touching, picking and manipulation. In addition to the development of the heterogeneous simulation environment, another contribution of this study was to demonstrate how lower-frequency threads, such as PhysX implicit integration, can fit in a much higher frequency interactive context, such as haptic rendering. Such unlimited combination of models, modes and techniques, although not trivial, make it possible to use the best solution available for each particular problem, opening up endless possibilities of what can be achieved with current hardware technology.

To demonstrate the effectiveness of our approach, we present the simulation of laparoscopic adjustable gastric banding (LAGB) as a case study. LAGB is a complex minimally invasive surgical procedure indicated for patients with morbid obesity. It consists of placing a flexible band around the stomach to constrict the food passage and produce an early satiety sensation. The motivation for this specific surgical procedure is the high cost in training surgeons and the increasing demand for it.

In this article we present an overview of the PhysX library, introduce surgery simulation with haptics and present our simulator. We then introduce a case study,

showing how we developed and modelled anatomy and physiology for the simulator, followed by a discussion of the results and conclusions.

Related work

Real-time physics-based surgical simulation is computationally very demanding (1); for a review of current literature on surgical simulation, please refer to (2). In summary, geometry-based models developed by, for example, Basdogan *et al.* (3,4), tend to sacrifice the physical realism of the deformation process to achieve real-time performance. Physics-based techniques that take into account the underlying mechanics of soft tissue deformation using mass spring networks (5) or finite elements (6,7) have been developed. Techniques such as condensation (8) and precomputation (9) have been used to accelerate finite element computations. Mesh-free methods, such as the point-associated finite field (PAFF), offer yet another alternative to finite element techniques and several methods have been presented to reduce computational costs (1). Hardware acceleration has also been used to achieve real-time rates (10).

As the capability of computing resources has increased over the years, so have the complexities of surgical scenes being simulated. A typical surgical simulation scene now involves much more than just tool–tissue interaction. Simulations of multiple organs, medical devices, sutures, coupled fluid flows due to bleeding and modelling of the physiological consequences of surgical procedures are now becoming common. To achieve a certain degree of realism in such complex surgical interactions, the simulation environments need to support the following: (a) heterogeneous scenes composed of different states of matter (solids, liquids and gases); (b) complex geometry and material properties of objects within the scene; (c) dynamic and real-time interaction (palpation, cutting, etc.) between virtual objects and tools physically manipulated by the user; and (d) multimodal (visual and haptic) rendering of the results to the user. It is clear that it is not possible to satisfy all four requirements using a single modelling technique. It is therefore necessary to use heterogeneous techniques, each optimized for one or more of the tasks. Real-time physics engines, originally developed for games development, offer a unique platform to develop such heterogeneous simulation scenarios.

In the literature, surgical simulators using physics engines are not very common. This happens mainly because the existing physical engines are often focused on a reduced set of physical laws, which is not enough for a full surgery simulation. Another reason is that they are implemented aiming either at producing fast, less accurate animations or very accurate but computationally demanding behaviours that cannot be produced in real time. However, a few examples of physics engines exist that are used for surgery simulation (11). However, these are not general physics engines for games. They are either

custom developments, in which the developers have the full control of the physical laws they are implementing, or general engines, allowing plugging in new physical laws using software containers. The former are not consolidated engines and demand great investment in development, while the latter are too general and cannot be optimized in terms of hardware use.

Physics engines that are generally targeted for games development are not specifically concerned with the intricacies involved in surgical simulation. Nevertheless, the problems that need to be tackled are common. For instance, the challenges involved in the interaction of an avatar and its clothing in a game scene is very similar to a surgical scene where the surgeon interacts with a membrane. Such similarities are explored in Figure 1.

Recently, several physics engines have been developed. Some of these are available commercially, have a closed source or are under limited licence, such as Havok (12), PhysX (13) and Newton Game Dynamics. Some other engines are open source and available free, such as ODE, Open tissue, Dynamechs, True Axis, Tokamak, Bullet, CMS-Labs Vortex, JIGLIB and BOX2D (14). For

more information about the simulator capabilities, we encourage the reader to look into the work of Seugling and Rölin (15) and Boeing *et al.* (16). They compared the available physics engines and described the features of each engine. One important conclusion is that only a few of them have the capability to simulate the surgical scenario.

In this regard, Marks *et al.* (17) investigated the game engines and looked into the adaptability of these engines to surgical simulation. They chose three game engines presenting functionalities that allow for the development of surgical simulations. Id Tech4 (18), Unreal Engine (19) and Source Engine (20) were chosen for further evaluation. These engines were evaluated in terms of editing, content and gameplay. More specifically for the surgical case, their assessment focuses on the maintainability of the surgical task. They analyse whether the engine supports change in surgical procedures, creation of surgical scenarios, simplified incorporation of other custom models, multiple users in simultaneous interactions and, finally, analyse the quality of physics-based interaction.

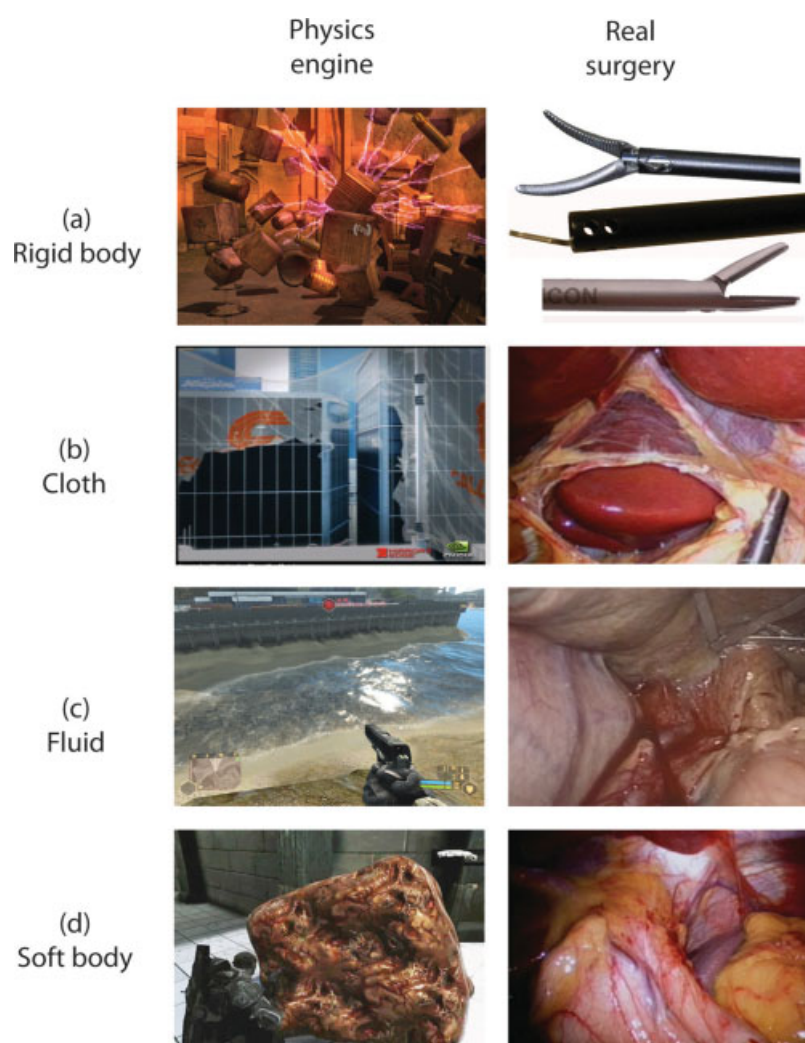


Figure 1. Entities available in PhysX match many surgical entities, such as: (a) rigid bodies may be used for modelling surgical instruments; (b) cloth for membranous tissue; (c) fluids for blood; and (d) soft bodies for organs

Physics-based interaction is at the intersection of physics-based engines and real-time haptic interaction. However, a good compromise between them is not easily established, as both require massive computation at different frequencies. In previous studies we dealt with the problem of rendering haptic feedback when very complex triangle meshes are involved (21,22). The problem is often solved at the level of collision detection, using optimized algorithms. Collision detection involves checking the geometry of the target objects for interpenetration; for detailed surveys on collision detection, see (23,24).

Computational methods

We first discuss techniques of employing the PhysX library to simulate various entities – soft bodies, rigid bodies and fluids – that are of interest in surgical simulation. We then point out specific issues related to multimodal surgical simulation involving haptics. One of the most critical problems is to be able to maintain various update rates for collision detection, physics-based interaction, visual and haptic rendering on the same model data structure. We present an extended version of a MVC framework to resolve this issue. We present concrete examples of how anatomy and surgical instruments may be modelled using various PhysX objects. Finally, we present our experiences in developing a procedural simulator for LAGB.

Using PhysX for interactive simulations

PhysX is a physics-based framework originally developed by Ageia, which was acquired in early 2008 by NVIDIA. Subsequently, NVIDIA added their GPU support in the PhysX framework to harness the capability of their general purpose computation framework CUDA (25). In general, the PhysX framework supports real-time simulation of various physical entities, such as rigid bodies, soft bodies, cloth, joints, fluids, springs and characters, as well as

collision detection and response algorithms, with a focus on application to the gaming industry. At the present time, the framework abstracts the hardware use from developers seamlessly and uses GPU or GPUs (if the system has multiple GPU cores) for the simulation without coding.

PhysX (13) is a simulation engine applying position-based dynamics (26). The framework provides functionalities revolving around the rudimentary physical objects provided, such as rigid bodies and joints, to more complicated physical objects, such as soft bodies, fluids and cloth (see Figure 2). PhysX mainly differs from game engines in the sense that it does not provide high-fidelity rendering and visualization. All rendering and visualization, and tasks related to network communication, user interactivity, audio and so on, is left to the developer.

In PhysX, entities are created as scene objects (NxScene). Scene objects are initialized by the NxPhysicsSDK. There can be more than one scene object at a time and each scene object is responsible for simulating the behaviour of the constituent objects.

The basic PhysX object is called an *actor*. There are two different actors in PhysX: static and dynamic. As the names imply, static actors are attached to a particular location and stay stationary most of the time in the scene. Conversely, dynamic actors are influenced by the forces in the environment and also influence other objects in the environment when collision happens. Collisions take place between the actors, whose shapes may vary from simple primitives, such as box, capsules, spheres, etc., to more complex geometries, such as convex mesh or triangle mesh.

The centre of mass of each actor is computed when the actor is created. Furthermore, the actor is moved and rotated based on the simulation dynamics referred to this centre. In the framework, actors have a body only if they are dynamic entities. The body is associated with *NxBodyDesc* class, which gives details regarding the body specification, such as mass, linear and angular velocity and local pose with respect to the centre of mass.

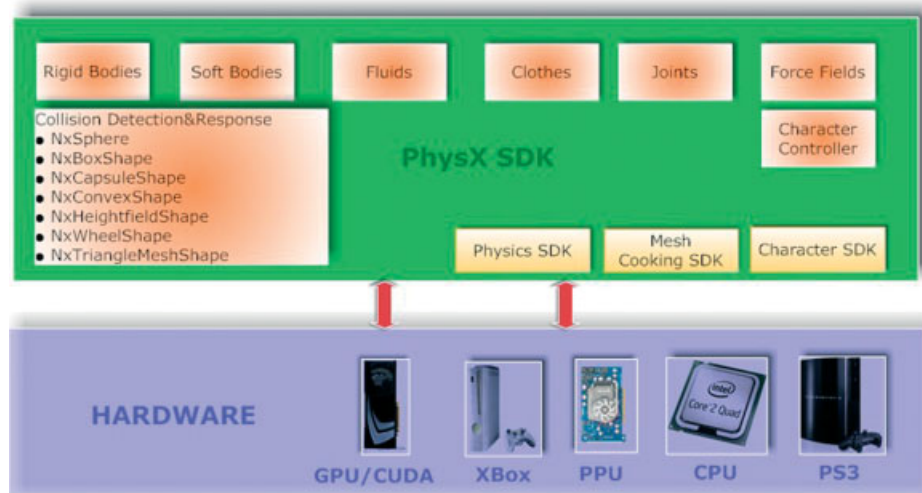


Figure 2. The PhysX architecture

The information also includes damping along the linear movement and angular movement. Actor body description also contains an inertia tensor that is calculated from its shape and density by the framework, or it can be explicitly set by computing the tensor manually. In the simulation, external forces, such as direct and impulse forces, can be applied to an actor body and are smoothed, so that the resultant force response effect can be controlled. Generally, in games, actors are primarily designed for rigid body elements.

PhysX provides *NxJoint* class to simulate joint structures. With joints, actors can be connected and motions can be limited within a desired range. There are several predefined joints available (spherical, revolute, prismatic). Apart from the existing joints, PhysX allows custom joints to be defined with the *NxD6Joint* class. This encompasses all existing joints and provides more functionality to create and freeze the degrees of freedom or limit the range of motion of each joint axis. The framework also provides a driver mechanism to simulate external joint movement.

PhysX fluids are essentially based on the particles approach. In this technique, the fluid is considered as particles that have density, velocity, position and lifetime states. The fluid is characterized by properties such as stiffness, damping, viscosity, etc. The particles interact with each other based on two different approaches, depending on user choice: a smoothed particle hydrodynamics (SPH) approach and a simple mode, where interparticle forces are not accounted for. Based on the desired fluid physics, *NxFluidDesc* identifies particle interaction. Currently, PhysX fluids can interact with PhysX cloth, soft bodies and rigid bodies. However, fluid–fluid interaction is not taken into consideration.

PhysX cloths consist of particles that are constrained within the edges of a shape using bending and stretching parameters. *NxClothDesc* offers additional parameters, such as damping, density, friction, pressure for closed cloths to simulate rather stiff behaviour (e.g. a rug) or more elastic behaviour (e.g. a skirt). In PhysX cloth, the collision detection takes place between a body and the cloth vertices. The collision detection sensitivity and the response to collisions can be manipulated by the thickness of the cloth and collision response coefficients, respectively. One of the remarkable characteristics of cloth is that it may be torn, which can be controlled with either user-supplied parameters that define the tearing factor or with explicit function calls.

Soft body simulation is one of the most notable features of the PhysX framework in the context of this work. The framework simulates any given topological elastic body. In general, it differentiates the simulation mesh geometry from the mesh that is used for rendering the soft body. Therefore, the SDK uses a tetrahedral mesh for the internal structure, which is encapsulated by the surface mesh of the body. The method of integration used is semi-implicit. Verlet integration is applied to compute the new positions of each tetrahedron node. The vertices' new positions under external and internal forces

(from springs connected to the tetrahedrons' vertices) are handled explicitly. In addition to this, the constraints solver is employed in the simulation loop, where the solver iterates over constraints by a predefined number of iterations which is explicitly set in the soft body description. This type of scheme, where both explicit and implicit integration methods are used at the same time, is called a semi-implicit scheme (26). In the constraint solver part, the non-linear Gauss–Seidel-type iteration is used where each constraint is handled individually. The solver actually resolves the collisions as well as predetermined constraints after the vertices are projected to their next time step positions by Verlet integration. Constraints solver both copes with fixed constraints and collision constraints. For instance, the stretching stiffness which defines a stretching constraint normalized between zero and one along the edges of a tetrahedron, or volume constraints that create a dependency between the initial reference volume and the current volume, can be regarded as fixed constraints. Thus, as stiffness and the number of iterations are increased, both volume and stretching results in stiffer behaviour and a slower frame rate. When the collision exists in a given step, the new vertex positions are updated with respect to the generated collision constraints. At the end of the simulation loop, the new positions comply with the fixed and collision constraints, which are then updated.

There are several parameters that influence the physics of the soft body significantly. Volume and stretching stiffness of the tetrahedrons specify the internal structure. On the other hand, damping and friction have more impact on the global response of the soft body. Naturally, the equation solver and the number of iterations have substantial effect on the softness and hardness of the body. Moreover, the SDK allows attaching rigid internal shapes to the bodies (a kind of skeleton) that can strengthen the body and provides more realism when the body has a skeleton in reality.

Surgery simulation with haptics

Force feedback (27) is an essential component of interactive simulations such as minimally-invasive surgery simulation (3), where the user interacts with soft, deformable organs using long, slender, surgical tools. The simplest paradigm of haptic interaction is the use of a point-based representation of the haptic cursor. However, this is unrealistic for many applications, including laparoscopic surgery simulation, where it may be necessary to retract an organ or appendage while grasping another one or to slice through tissue, fat or fascia using blunt or cautery instruments. For such applications, it is natural to represent the haptic cursor as a line or geometry. However, for deformable organ models, such ray- and geometry-based representations are prohibitively expensive if realistic organ models are used, and realistic interaction is necessary.

There are many existing interaction approaches that are aimed at reaching visually interactive speeds, i.e. around 30 frames per second, for real-time applications. However, in virtual environments involving haptics, a much higher update rate of at least a few hundred frames per second is necessary to render stable force feedback. If physics-based computations are to be performed as part of the model behaviour, then this places severe demands on the efficiency of the algorithms.

We have developed a framework based on the MVC pattern (28), which allows for the computation, visualization and interaction with physics-based models. MVC is a design and architectural pattern used in software engineering. It aims at isolating data, business logic and how information is displayed. As a result, changes in the visual appearance of the application are independent of the business rules and vice versa. In MVC, the *model* represents the application data; the *view* corresponds to elements of the user interface, such as text, widgets, 3D rendering areas and so on; and the *controller* manages the data manipulation and modification applying the business rules. The controller reads data from the model, changes them and writes them back to the model, so that they are available to be viewed or modified by other controllers. Many controllers and/or viewers are allowed simultaneously, while all data are seen as a unique module. In our framework, *data* hold geometric and material information, *view*s display *model* information on graphics windows and user interface elements such as key pressing or mouse clicking, and *controllers* modify the *model*, e.g. the collision detection controller inspects the geometry and flags the penetrating triangles, the physics controller applies physical rules to modify the forces, accelerations, velocities and positions in the *model*, and so on. While similar control schemes have been used before (26), they usually prioritize the physics simulation over other control tasks, keeping it apart from the MVC structure. The uniqueness of our approach is that physics is also encapsulated in a controller, managed by PhysX, resulting in a highly efficient use of the hardware resources.

The haptics controller allows communication of the model with the force feedback device (the phantom). This controller is not responsible for calculating the forces. The interaction forces are calculated by the physics and collision detection controllers and stored in the model. Hence, the haptics controller becomes a very fast scanner which reads forces and positions from the model and makes the matrix conversions necessary before calling back the device driver. Such conversions are often necessary to register 3D spaces and to apply scaling between worlds. Since we have used the Sensable Phantom Omni™ as force feedback device, the OpenHaptics library is the interface used by this module to communicate with the device. Such asynchronous design allows the force feedback to be updated at a frequency of more than 10 kHz. Notice that it is a consensus in the haptics literature that at least 1 kHz is necessary to ensure smooth transitions.

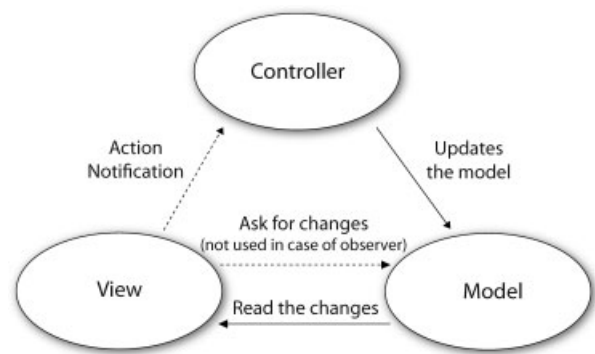


Figure 3. The model-view-controller (MVC) pattern

Our framework is designed in such a way that multiple execution threads run in parallel and concurrently at different rates and in different cores on the latest CPUs. This permits, for example, that haptic computation is transparently made at over 1000 Hz, while a custom explicit integrated mass-spring system runs at lower frequencies, below 400 Hz. At the same time, graphics rendering is made at 30 Hz. This is possible because with MVC only the controllers modify the model (see Figure 3). While this approach allows simulating some surgical steps, it cannot completely simulate a surgery because, due to computational complexity, only a few entities can be physically based and increasing the number of entities degrades the performance dramatically. To keep the performance at the high level required for the task, we replaced the custom explicitly integrated mass-spring model by PhysX, which is implicitly integrated and thus much more stable. The trade-off is that implicit integration methods are more computationally demanding. With our current soft bodies, the physics thread runs at below 20 Hz. At the same time, collision detection, which is optimal at a much higher frequency depending on the dynamics of the simulation (objects moving fast may require over 1000 Hz, while calm scenes would accept collisions to run at the same pace as physics) and realistic graphical rendering, which runs at a constant frequency of 30 Hz, must be maintained. For this to happen, the model data structure, viewer and interaction controllers must remain unchanged, which is not straightforward.

We solve this critical problem by developing an extended version of the MVC framework, as presented in Figure 4. In this framework, we include PhysX as a new controller that modifies the model geometry already present in the data structure. Every PhysX step is performed on a separate thread at its optimal frequency, determined by the physical properties of the materials involved. At the end of each step, the model geometry is updated at once, similarly to a swap buffer step.

Modelling anatomy, physiology and tools

For the assessment of our simulator, we considered a realistic scenario of a laparoscopic surgical procedure.

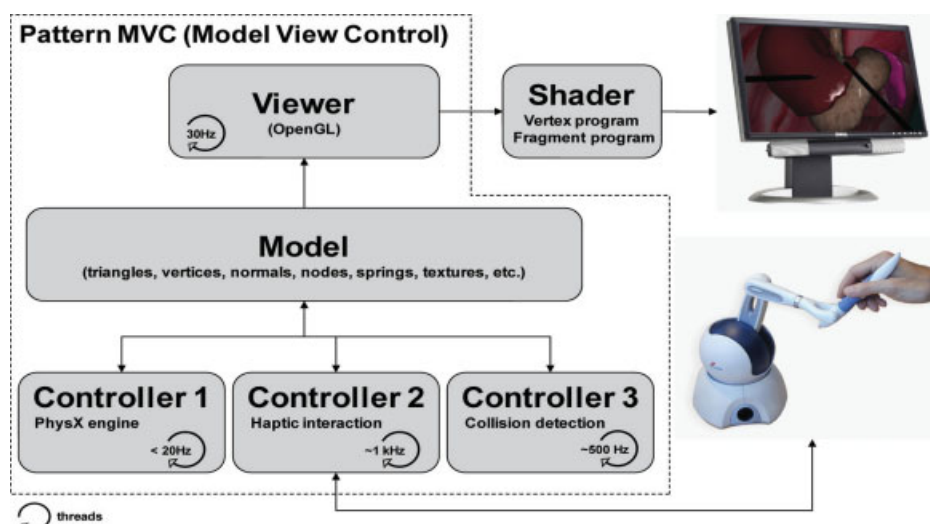


Figure 4. The extended MVC framework proposed to deal with different update rates required for collision detection, physics-based interaction, visual and haptic rendering on the same model data structure for surgery simulation

One of the goals of this experiment was to evaluate our algorithms in the practical situation of a virtual laparoscopic intervention and to explore the responsiveness of the system. An additional aim was to demonstrate that high frequency haptic response and rigid and deformable objects simulated at a lower frequency by implicit integration can co-exist in a fully interactive real-time haptic environment.

The scenario consisted of a partial model of the interior of the abdominal cavity, where a few organs were visible, as they would be visible through the laparoscope. The organs included a deformable liver, a deformable stomach and a rigid spleen, with the peritoneum in the background.

The organ models were obtained from the segmented data of the VIP-Man, which is a series of high-resolution photographic images from the Visible Human Project dataset (29). Images we used were 1878 slices of 1760×1024 pixels in raw format. The pixel size was 8 bits, with a value range of 0–71 according to the labels in Table 1. In the original images, the distance between slices was 3 mm and between pixels it was 0.35 mm; however, we down-sampled them to pixels of 1.4 mm for performance issues; see a slice example in Figure 5. As the stomach was labelled in three different parts – wall, content and mucosa – we relabelled all three to the same value. The oesophagus was also blended with the stomach to avoid non-realistic discontinuities. Then we used the VTK (30) implementation of the marching cubes algorithm (31) to extract the organs' surfaces. As the marching cubes algorithm is based on an isosurface, we selected the tissue to be reconstructed setting image values in the range of the targeted tissues to 1024, while all others were set to 0. An isosurface was computed with the value 512, marching cubes placed triangles at the cube–surface intersections and the resulting dataset was filtered to reduce the number of triangles before being written to a file in standard formats such as obj and stl. The liver, stomach and spleen were meshes composed of

2484, 1863 and 1040 triangles, respectively. The liver and stomach were simulated by PhysX soft bodies with around 3 k tetrahedrons.

The peritoneum is the internal wall of the abdominal cavity and was not labelled in the VIP-Man. Hence, a different method was used to create the mesh. We drew outlines (Figure 6) on selected slices of the cross-sectional

Table 1. Some tissues and their corresponding labels in the VIP-Man image files

Organ	Label
Oesophagus	12
Oesophagus lumen	58
Oesophagus mucosa	59
Fat	29
Liver	35
Spleen	40
Stomach wall	33
Stomach content	48
Stomach mucosa	52




Figure 5. Example slice from the VIP-Man dataset




Figure 6. Outline of the peritoneum




Figure 7. Schematic of the peritoneum reconstruction from image slices. Outlines on each pair of consecutive slices have their nodes connected to form triangles

Visible Human CT dataset, composed of a series of points which were dense in curved areas and sparse in areas of lesser curvature. We started with the most distal cross-sectional layer, manually outlined the shape on the image, then copied the outline and dragged the copy along the proximal direction to the next slice. Control points were edited (moved/added/united) on the new slice to outline the peritoneum shape according to the corresponding anatomy. Triangles were created to mesh the space between the two outline curves. The process was repeated for all subsequent layers until the peritoneum was no longer visible for the current layer. Figure 7 shows a schematic of the process.

The resulting peritoneum tightly enclosed the organs. However, for surgery, the abdomen is insufflated with CO₂ to provide room for the laparoscope and other tools. To simulate the shape when the peritoneum was insufflated, we edited the control points to enlarge the front peritoneum. The resultant peritoneum, as shown in Figure 8a, contained 213 vertices and 384 triangles. We further used texture mapping and per-pixel illumination models to render realistic graphics for all organ models, as shown in Figure 8b. As we will see in the next section, the surgical scenario we simulated contained a band known as the LapBand[®], developed by Allergan Inc. (see Figure 9a), that was placed around the stomach model.

Modelling of this band is not trivial. The LapBand[®] has two distinct parts having different physical behaviour. The thick and stiffer part, which is identified as the upper part, has an adjustable ring. The long part, which looks like a tube, is more flexible, threads through the stiffer part, and connects the band to the port in the actual surgery. The thick part is almost inextensible and tends to return to its original curved shape when released. The long part is more elastic and highly deformable. Since the two regions have different material and structural properties, we split the band model into two different meshes so that it presented plausible physical behaviour. It is worthwhile to mention here that a limitation of PhysX is that soft bodies must be isotropic and homogeneous, i.e. it does not allow the definition of two different material properties within the same soft body. Having two distinct meshes gave us the flexibility to define physical parameters independently between them. For instance, stretching stiffness was set to the maximum value for the thick part. In addition, the thick part had a high factor of volume conservation, was highly damped, and the global damping (NX_SBF_COMDAMPING) was also highly effective in its deformation. On the other hand, the long part was less stiff and no global damping was used (centre of mass damping).

Since the two meshes were detached, some sort of constraint had to be established between them to avoid discontinuous motion when either of the parts was moved or deformed. The solution we proposed using PhysX was to attach a rigid body actor to either parts of the band and then connect them with a distance joint. The distance joint provided constrained distance in each iteration and allowed for the motion to be guided through any one of the two meshes. The solver number of iterations per second for the band was set to a higher value, which affected the stability of the overall simulation.

The haptic cursor on the screen may take the shape of any surgical instrument. When an instrument is selected, it is associated with the phantom cursor and moves following motion of the handle manipulated by the user. The user can switch between instruments at any time using a menu. Currently, we modelled the set of instruments required for the LAGB procedure. This included the cautery hook, the grasper and the scissors. The cautery hook was modelled as a unique rigid body on PhysX. The grasper and scissors were each modelled as a set of three PhysX convex rigid actors. Both had one root part, which moved following the phantom stylus, and two leaf parts, which performed an open/close motion to cut in the case of the scissors, and to grasp in the case of the grasper. All three parts were rigid bodies in PhysX and were connected by a rotational joint. Figure 10 illustrates the instrument models.

Laparoscopic adjustable gastric banding (LAGB) procedural simulation

We used the procedural simulation of LAGB as a case study. LAGB is a complex minimally invasive surgical




Figure 8. Bottom view of the peritoneum model and interior organs in wireframe (a) and textured (b)




Figure 9. (a) The Allergan LAP-BAND® system and (b) our model. The thick part of the band is stiffer than the long part and the two parts are connected by a distance joint

procedure which consists of placing a Lap Band® around the stomach of a morbidly obese patient to constrict the passage of food, which results in early satiation. We simulated the 'pars flaccida' technique for the lap-band placement, which basically involves three phases:

Phase 1: port placement

In this phase, five access ports are placed on the abdomen. These are small perforations in the abdominal walls through which the trocars and laparoscope will be introduced. This is followed by pumping carbon dioxide

gas into the peritoneal cavity (pneumoperitoneum), which separates the organs and allows better accessibility.

Phase 2: Dissection

In this phase a passage is created behind the upper part of the stomach to pass the band. To perform this step, the liver is retracted using a Nathanson hook liver retractor, the gastro-hepatic ligament (pars flaccida) is divided, avoiding the hepatic artery, followed by retraction of the gastric fundus and omental fat. The angle of His is subsequently mobilized to create a small window in the phreno-oesophageal ligament. The right crus is identified and the peritoneum overlying its lower portion near the junction of the left crus is divided.

Phase 3: Band placement

In the next phase, a grasper is inserted from the right, passing through the opening behind the upper part of the stomach until it can be seen by the angle of His at the left. The band is then inserted into the abdomen and pulled around posteriorly, from left to right, with the passing device or a grasper. The band is then secured and locked in position.

As the first phase is extracorporeal, we focus here on the next two phases. The major steps in this modelling using PhysX are explained below, with Figure 11 illustrating these steps:

Modelling phase 2

We modelled a hook cautery (Figure 10a) as a rigid body and a blunt dissector or grasper (Figure 10b) as a rigid articulated body. The positions and orientations of the surgical instruments (hook cautery, blunt dissector or grasper) were controlled using a Phantom Omni as 3D input device, with six degrees of freedom (DOF) for each instrument. When the virtual instruments contact a virtual organ, the haptic interface device produces force feedback to the user's hand. The electrosurgical




Figure 10. Models of the laparoscopic instruments: (a) hook cautery; (b) grasper; (c) scissors

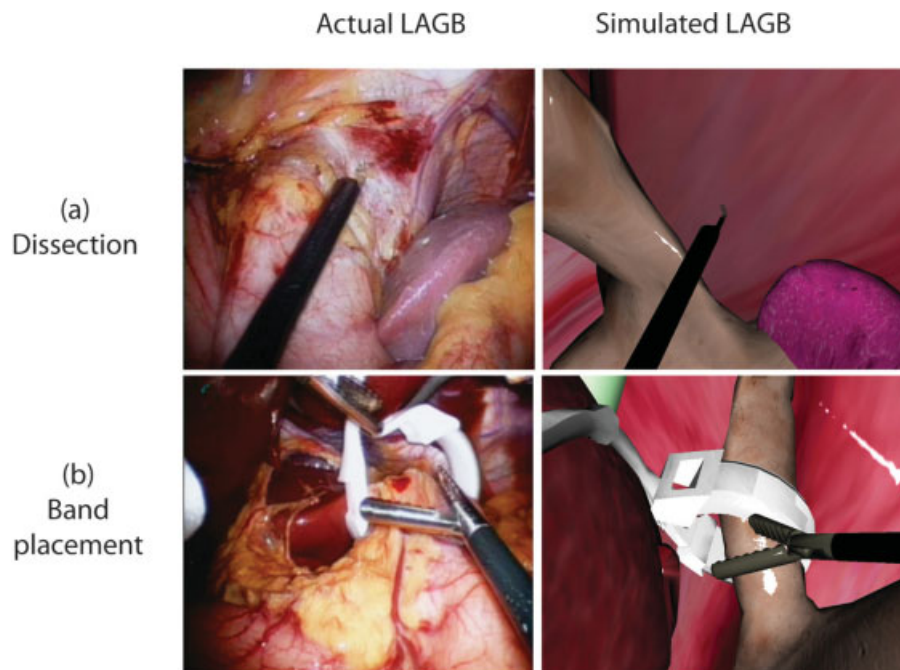


Figure 11. Relating actual surgery video frames with screenshots of the simulator: (a) tissue dissection (phase 2); (b) band placement (phase 3). Not all anatomical structures are shown in the current simulated LAGB. Further developments will include fat, ligaments and other important anatomical landmarks




Figure 12. Electrosurgery is used to cut the tissue. (a) Colour mapping shows the temperature distribution; (b) result of cutting a membranous tissue; (c) result of cautery applied to an organ

procedure is simulated by modelling the temperature increase caused by approaching the cautery to the tissue, removing triangles when vaporization temperature is reached (32). Figure 12 shows the colour mapping of the temperature (a), the result of cutting a membranous tissue (b) and the result of the cautery applied to a massive organ (c). The dissector can grab parts of the organs to push and pull soft regions in order to remove obstacles from the field of view, allowing visual exploration of

hidden areas, such as the right crus and the angle of His. The virtual laparoscope is controlled with the mouse.

To accomplish these tasks, we combined custom methods and PhysX capabilities. The organ deformation was calculated by PhysX, while the graphics rendering was accomplished using our custom shaders. As soft body to soft body collision detection is not efficient in PhysX, the method described in reference (33) was used for organ to organ collision detection. However, PhysX rigid

body to soft body collision detection is very efficient and was adapted to detect contacts between instruments and organs.

Modelling phase 3

The instruments used in phase 2 were also used in this phase. Two graspers were used to manipulate the band. As discussed in the previous section, the band was modelled as two meshes which were coupled and could be manipulated (picked and dragged) with the instruments. Consequently, collision detections between the band and the organs (soft bodies) required a custom method (33), but collisions with the instruments were dealt with well by PhysX. As for picking, when the grasper closed, all colliding triangles were flagged as *picked* by that instrument. Then, as the closed grasper moved around, all picked triangles were kinematically controlled in the PhysX model, i.e. their positions were updated following the displacements of the handle at the beginning of every simulation frame. Finally, the positions of all nodes adjacent to flagged triangles were recalculated by PhysX to smoothly follow the displacements. Flags were reset when the grasper opened to release the picked triangles.

Results

We have implemented an interactive PC-based surgical simulation framework and tested it on an Intel(R) Core™ 2 Quad 2.66 GHz machine with a GeForce 8800 GTX graphics card. Customized vertex and pixel programs were used for textured shading, which provided state-of-the-art interactive graphics realism. This simulator utilized two force feedback devices to provide bimanual interactivity. The hardware set-up is shown in Figure 13.

The framework relies on a combination of the PhysX physics engine and customized parallel algorithms to render physics, graphics and haptics in real-time. It exploits the parallel capabilities of current multi-core

CPUs and multiprocessor GPUs to maximize efficiency. Although using game-oriented engines has the potential to improve efficiency and physics realism, the greatest challenge was to control the various threads running at different frequencies in the same environment. We showed that this is possible by extending the MVC pattern, in the sense that the many *controllers* exchanged information through the *model* at the frequency of the slower thread for each pair of threads.

The system displays graphics at 60 Hz, which, for stereo visualization, is 30 Hz for each monitor/eye. Such frequencies provide smooth graphical displays with no flickering. The haptic response is provided at frequencies over 1000 Hz, which allows for a vibration-free haptics rendering. The force update in the model, in turn, is bounded by the collision detection and response algorithms, which run at frequencies of 100–4000 Hz, depending on the number of elements involved: rigid to rigid being the fastest; soft to rigid being slower; and soft to soft being the slowest.

As for the physics processing, while explicit integration methods are simple but require more steps (usually 300 to 600 steps per second are explicitly set for real time with our models) for convergence, PhysX uses implicit methods that are more time consuming *per frame* but are unconditionally stable. Typically, our simulation processes physics at 10–20 Hz. It is important to notice that at times between the 10–20 Hz updates the tools can still be moving fast (updated at several hundreds of Hertz) and then, even if the soft body vertex positions are not updated in that interval, the force feedback resulting from collisions will change at collision frequency as the tools move. Regarding physics implementations on GPU and CPU, in preliminary tests we observed that some models run over 10 times faster when the parallel hardware (GPU) is enabled, especially when fluid simulations are involved. However, other models run at the same rate with both GPU and CPU implementations. This happens because the current release of the PhysX driver enables only a small set of its functions to run on the GPU through the CUDA technology. Consequently, to make any conclusion here in this regard is premature. We plan to perform thorough comparative tests with different scenarios as soon as a new driver is released, which is expected this year.

Another issue is deformation accuracy. To demonstrate that PhysX soft bodies are accurate enough for surgery simulation, we simulated the hemisphere of Figure 14a. We modelled the same hemisphere with both PhysX and finite elements (FEM). For the PhysX model we used 1258 vertices, 3901 tetrahedra and 4224 links, and for the FEM discretization we used 1715 nodal points. Then we indented the two hemispheres at their poles to simulate the interaction with a surgical tool tip and analysed the resulting deformation profiles. Results in Figure 14b show that the profile obtained with PhysX was close enough to that obtained with FEM for the same indentation of the pole.



Figure 13. Hardware set-up of the simulator




Figure 14. The hemisphere problem. (a) The hemisphere as modelled in PhysX with 3901 tetrahedra and the applied indentation. (b) The resulting profiles for PhysX and FEM deformation and the undeformed profile




Figure 15. Normalized relations between some physical parameters. As a higher stiffness causes the numerical integration to diverge, higher densities make it more stable. However, increasing the number of time steps improves stability but increases computational time

We further analysed how execution time and stability change with the number of polygons and PhysX parameters, such as stiffness, density, damping, etc., but especially how scale changes the physical behaviour. We noticed that scaling the original model affects the responsiveness of PhysX, due to the higher stability of massive elements in the process of numerical integration of physical behaviour. Stiffer models, as expected, require additional steps in numerical integration and add delay to the physics response. Similar behaviour occurs when the density of the elements is too low. Damping the system increases stability but also increases delays. Due to this two-way relationship between density, damping and stiffness, as shown in Figure 15, if changes in the scale of the model geometry are distributed in the same proportions to the mechanical parameters, then it will cause behavioural disturbances. The scale factors for each property must be defined taking adequate proportions into account. Further details on scalability are discussed in (34).

Discussion

The complete LAGB simulation is complex and under development. However, important advances have been made in developing the overall simulation framework based on PhysX. There are many advantages of using a

physics engine such as PhysX, as outlined above. The results are encouraging and translate easily to other surgical simulation scenarios with complex interacting soft bodies. However, it is important to realize some of the limitations of PhysX, such as the inability to describe multiple material properties of the same object and the lack of a mechanism to modify accelerations and velocities during a simulation step. The PhysX source code is closed and the API does not allow integration of custom algorithms. Finally, the collision detection algorithms do not cover all primitive to primitive tests.

It is important to realize that surgical simulators are not glorified video games. Soft tissue mechanical characteristics must be obtained from actual experimental data. There is significant work being conducted in our laboratory in this direction (35). Another important aspect is validation. The simulators must be tested by surgeons and surgical residents to ensure that they indeed result in positive training transfer. In fact, to justify their added costs, they must be more effective than traditional mechanical training tool boxes or video trainers. Such validation efforts are also under way at Harvard Medical School.

Acknowledgements

Partial support of this work through Grant No. R01 EB005807-01 from NIH/NIBIB and Grant No. 151820/2007-4 from CNPq, Brazil, is gratefully acknowledged.

References

- De S, Lim Y-J, Muniyandi M, *et al.* Physically realistic virtual surgery using the point-associated finite field (PAFF) approach. *Presence* 2006; **15**: 294–308.
- Lim Y-J, Jin W, De S. On some recent advances in multimodal surgery simulation: a hybrid approach to surgical cutting and the use of video images for enhanced realism. *Presence* 2007; **16**: 563–583.
- Basdogan C, Ho C-H, Srinivasan MA, *et al.* Force interactions in laparoscopic simulations: haptic rendering of soft tissues. *Medicine Meets Virtual Reality Conference 6*, San Diego, CA, 1998.
- Delp SL, Loan JP, Basdogan C, *et al.* Surgical simulation: an emerging technology for emergency medical training. *Presence* 1997; **6**: 147–159.
- Kühnapfel UG, Neisius B, Krum HG, *et al.* CAD-based simulation and modelling for endoscopic surgery. *Proceedings of MedTech*, Berlin, Germany, 1994.
- Bro-Nielsen M. Finite element modeling in surgery simulation. *Proc IEEE* 1998; **86**: 490–503.
- De S, Srinivasan MA. Thin walled models for haptic and graphical rendering of soft tissues in surgical simulations. In *Medicine Meets Virtual Reality 7 (MMVR7)*. IOS Press: Amsterdam, 1999; 94–99.
- Bro-Nielsen M, Cotin S. Real-time volumetric deformable models for surgery simulation using finite elements and condensation. *Computer Graphics Forum* 1996; **15**: 57–66.
- Picinbono G, Delingette H, Ayache N. Non-linear anisotropic elasticity for real-time surgery simulation. *Graph Models* 2003; **65**: 305–321.
- Liu Y, De S, Westwood JD, *et al.* (eds). CUDA-based real-time surgery simulation. In *Medicine Meets Virtual Reality 16 (MMVR16)*. IOS Press, Amsterdam, 2008; **132**: 275–280.
- Nourian S, Xiaojun S, Georganas ND. Role of extensible physics engine in surgery simulations. In *IEEE International Workshop on Haptic Audio Visual Environments and their Applications*, Ottawa, Ontario, Canada, October 2005; 6 pp.
- Havok Havok Home Page: www.havok.com, 2008.
- NVIDIA PhysX library: www.nvidia.com/object/nvidia_physx.html, 2008.
- Box2D.org Box2D Physics Engine: box.org, 2008.
- Seugling A, Rölin M. Evaluation of physics engines and implementation of a physics module in a 3D-authoring tool. Masters Thesis, Department of Computing Science, Umea University, 2006.
- Boeing A, Bräunl T. Evaluation of real-time physics simulation systems. In *GRAPHITE '07: Proceedings of the 5th International Conference on Computer Graphics and Interactive Techniques in Australia and Southeast Asia*. ACM: New York, 2007; 281–288.
- Marks S, Windsor J, Wünsche B. Evaluation of game engines for simulated surgical training. In *GRAPHITE '07: Proceedings of the 5th International Conference on Computer Graphics and Interactive Techniques in Australia and Southeast Asia*. ACM: New York, 2007; 273–280.
- Wikipedia. ID Tech 4 (Doom 3 engine): en.wikipedia.org/wiki/Id.Tech.4, 2008.
- Unreal. The unreal engine: www.unrealtechnology.com, 2008.
- Valve. The source engine: source.valvesoftware.com, 2008.
- Maciel A, De S. An efficient dynamic point algorithm for line-based collision detection in real-time virtual environments involving haptics. *Comput Anim Virtual Worlds* 2008; **19**(2): 151–163.
- Maciel A, Liu Y, Ahn W, *et al.* Development of the VBLaST™: a virtual basic laparoscopic skill trainer. *Int J Med Robotics Comput Assist Surg* 2008; **4**(2): 131–138.
- Lin MC, Gottschalk S. Collision detection between geometric models: a survey. In *Proceedings of the (IMA) Conference on the Mathematics of Surfaces*, Birmingham, UK, 1998; 37–56.
- Teschner M, Kimmerle S, Heidelberger B, *et al.* State of the art report: collision detection for deformable objects. In *Computer Graphics Forum*, Vol. 24, no. 1, pp. 61–81, March 2005. Blackwell Ltd. 2005, 2005.
- NVIDIA Compute unified device architecture (CUDA) framework. www.nvidia.com/object/cuda_home.html, 2008.
- Müller M, Heidelberger B, Hennix M, *et al.* Position-based dynamics. *J Vis Comun Image Represent* 2007; **18**(2): 109–118.
- Maass H, Chantier BBA, Çakmak HK, *et al.* How to add force feedback to a surgery simulator. In *Proceedings of the International Symposium on Surgery Simulation and Soft Tissue Modeling*, Ayache N, Delingette H (eds), 12–13 June 2003; 165–174.
- Gamma E, Helm R, Johnson R, *et al.* *Design Patterns: Elements of Reusable Object-oriented Software*. Addison-Wesley Professional: London, UK, 1995.
- Spitzer V, Ackerman MJ, Scherzinger AL, *et al.* The visible human male: a technical report. *J Am Med Inform Assoc* 1996; **3**: 118–130.
- Schroeder W, Martin K, Lorensen B. *The Visualization Toolkit: An Object-Oriented Approach to 3-D Graphics*. Prentice Hall: New Jersey, NJ, 1997.
- Lorensen WE, Cline HE. Marching cubes: a high resolution 3D surface construction algorithm. In *Proceedings of SIGGRAPH 87*, Anaheim, California, 1987; 163–169.
- Maciel A, De S. Physics-based real-time laparoscopic electrosurgery simulation. *Studies in health technology and informatics*. In *Medicine Meets Virtual Reality 16 – Parallel, Combinatorial, Convergent: NextMed by Design*, Westwood JD, *et al.* (eds). IOS Press: Amsterdam, 2008; **132**: 272–274.
- Maciel A, Boulic R, Thalmann D. Efficient Collision Detection within Deforming Spherical Sliding Contact. *IEEE Trans Visualiz Comput Graphics* 2007; **13**: 518–529.
- Barenblatt GI. *Scaling, Self-similarity, and Intermediate Asymptotics: Dimensional Analysis and Intermediate Asymptotics*. Cambridge University Press: Cambridge, 1996.
- Lim YJ, Deo D, Singh TP, *et al.* *In situ* measurement and modeling of biomechanical response of human cadaveric soft tissues for use in physics-based laparoscopic surgical simulation. *Surg Endosc* (E-pub ahead of print: DOI: 10.1007/S00464-008-0154-Z).